Continual Harness Policy: Reset-Free Online Adaptation

1. Architectural Mandate

The **Continual Harness** implements reset-free online adaptation for autonomous foundation agents operating within #d.u.m.b.a.s.s. It enables agents to improve from real-world execution feedback (such as syntax mistakes, model timeouts, and operator corrections) without suffering catastrophic forgetting or destabilizing production prompt stacks.

---

2. Core Operational Rules

Rule 1: Immutable Base Anchor

- Base system prompts (`system.md`) and core behavioral invariants are strictly read-only during runtime.
- Agents are **forbidden** from directly modifying their base prompt files.

Rule 2: Layered Prompt Refinement (`/refine`)

- Online adaptations are expressed strictly as **delta layers** appended to the agent's contextual skills:

  [Immutable Base Prompt] + [Domain Rules] + [Layered Refinement Delta]

- Deltas are versioned with semantic commit tags (`refine_v1.0.1`, `refine_v1.0.2`).

Rule 3: Automated Snapshot Rollbacks

- Before any delta layer is activated, the harness takes a cryptographic snapshot of the current state (`snapshot_hash`).
- The agent must pass a standardized regression test suite (Car Wash test harness) before the delta is marked permanent.
- If regression is detected (regression score drop > 2%), the harness automatically executes an instant rollback to the previous snapshot without requiring operator intervention.

---

3. Feedback Loop & Verification

1. **Trigger**: An error occurs during tool execution or the operator provides an inline correction.
2. **Analysis**: The triage auditor identifies the root cause and drafts a candidate refinement rule.
3. **Verification**: The candidate rule is tested against historical test cases in the Reasoning Bank.
4. **Adoption**: Upon passing verification (+1.0), the refinement rule is written to the agent's local `skills/` directory and synced to `#d.u.m.b.a.s.s.`.