




what's the recursive training through verification
model that Frontier platforms use for training their AI

August 8, 2026

11 sites

Frontier AI labs use recursive training loops driven by external verifiers (often part of **Recursive Self-Improvement** or **RSI** frameworks). A model generates problem-solving paths or code, and an independent symbolic, mathematical, or programmatic verifier checks the correctness. Validated reasoning traces are then fed back to train successive model generations.



Key Components of the Loop

- **The Proposer/Generator:** An AI agent that creates code, math proofs, or logic steps.
- **The Exogenous Verifier:** An external compiler, execution sandbox, or math checker that deterministically confirms if the output works.
- **The Feedback Filter:** A mechanism that strips out flawed reasoning (even if it yielded a lucky correct answer) to prevent model collapse.
- **Recursive Retraining:** Using verified synthetic data pairs to tune the next-generation model via reinforcement learning.

Why Labs Rely on It

- **Data Scarcity:** Overcomes the ceiling of human-generated text on the open web.
- **Scalable Reasoning:** Expands capability in math, coding, and logic where ground-truth verification is computationally free.
- **Efficiency:** Minimizes the need for continuous raw scale by optimizing reasoning performance via structured verification.

Further Exploration

- Read a technical breakdown of model self-review risks in the [arXiv Paper on Code LLM Collapse](#).
- Explore the mechanics of verifiable math training in the [Verify-RL Paper](#).
- Review industry discourse on recursive feedback loops from [Anthropic's Institute](#).

Would you like to explore how **exogenous verification prevents model collapse**, or look closer at **inference-time scaling versus training-time recursion**?



arXiv

Stabilizing Iterative Self-Training with Verified...

Contributions (4) We introduce NSPCA

Medium

The Verifier Bottleneck - by...

Only verifiers add

arXiv

When AI Reviews Its Own Code: Recursive Self-Trainin...

We therefore compare three recursive

Communications of t

Is Recursive Self-Improvement Real.

On February 5, 2024

You are completely right, my mistake—the four-letter acronym used by Frontier platforms is **RLVR**, which stands for **Reinforcement Learning with Verifiable Rewards**. YouTube · Vinh ... +1

Instead of relying on human judges or tricky, hallucination-prone AI reward models, **RLVR** hooks the training loop up directly to hard, objective, programmatic verifiers.

How the RLVR Framework Works

- **The Objective Task:** The model is fed problems with clear, deterministic "right or wrong" answers, like software code, complex math equations, or strict schema validation. GitHub +1
- **The Sandboxed Verifier:** When the model outputs a solution, an external, non-neural system (like a Python code compiler, code interpreter, or automated math grader) executes it. Level Up Co... +1
- **Binary/Verifiable Rewards:** The model receives a crystal-clear, 100% accurate binary score (e.g., 1.0 for passing all unit tests, 0.0 for failing).
- **Search Compression:** This exact feedback doesn't just expand what the model knows; it trains it to find the correct reasoning path on the *first* try instead of taking eight random guesses, compressing its search efficiency.

Why RLVR replaces traditional alignment methods

While standard Reinforcement Learning from Human Feedback ([RLHF](#)) teaches a model *style* and *safety*, **RLVR** is what teaches it *procedural reasoning* and true problem-solving. It completely bypasses the risk of "reward hacking" because the model cannot lie or sweet-talk its way into a high score with a compiler. GitHub +2

If you are interested, we can look into **how labs scale RLVR compute at test-time**, or dive into **writing a custom programmatic verifier wrapper**. Which sounds better to you? YouTube · AI Ma... +1



GitHub
RLHF vs RLAIF vs RLVR: The Three...
You can also combine methods. Many...

GitHub
Awesome RLVR — Reinforcement...
* Key: Uses binary, verifiable reward functio...

YouTube · Vinh Nguyen
Reinforcement Learning with...
all right today we're going to talk about something...

YouTube · AI Makers
RLVR: Reinforcement...
all right Whiz today we ta RL with verifiable reward.

When you mix **RLVR** (Reinforcement Learning with Verifiable Rewards) with **KAG (Knowledge-Augmented Generation)**, you get an AI model that doesn't just look up facts—it actively reasons through massive enterprise knowledge bases with strict, mathematically verifiable constraints.

How RLVR and KAG Fuse Together

The mixture occurs by turning KAG’s logical operations into an **RLVR environment**. Instead of checking just a simple string output, the external verifier checks the step-by-step logic paths the model takes through the graph.  [GitHub](#)

1. Formulating the "Logical Form" Prompt

When a user asks a complex, multi-hop question (e.g., *"Which of our Q3 clients in Europe experienced a supply chain delay matching the 2025 compliance schema?"*), the **KAG Solver** translates the question into a structured, step-by-step "Logical Form" execution plan.  [YouTube · Anil Nat...](#)

2. Model Trial Generation (The RLVR Step)

The LLM acts as an agent and generates several candidate execution traces (actions) to navigate the Knowledge Graph, fetch text chunks, and perform data calculations. 

3. Programmatic Graph Verification

This is where **RLVR** takes over from normal generation. An external graph engine (like the [OpenSPG Engine](#) used in foundational KAG frameworks) acts as the deterministic verifier. It runs the model's generated plan and checks for hard-coded compliance: 

- Did the entity alignment match a verified schema?
- Did the multi-hop path lead to a logically sound conclusion without broken links?
- Are the math and code outputs derived from the facts fully executing? 

4. Verifiable Reward Updates

If the reasoning trail successfully compiles and resolves against the Knowledge Graph, the path receives a full binary reward (1.0). If it hits a dead end, hallucinates a relationship, or breaks an entity constraint, it gets a 0.0.  [Promptfoo +1](#)

The Massive Benefits of Mixing RLVR + KAG



- **Zero-Hallucination Multi-Hop Reasoning:** Standard RAG breaks down when an answer requires joining 3 or 4 disparate documents together. By training the model via RLVR directly on KAG structures, the model learns the exact mechanics of entity normalization and graph navigation.  [Medium +1](#)
- **True Traceability:** Because RLVR rewards paths that ground themselves strictly back to the source text chunks via mutual indexing, the model is systematically penalized for trying to bypass the database to guess an answer from its pre-trained memory.  [YouTube +1](#)
- **Emergent Agentic Troubleshooting:** When trained in this environment, models develop the ability to self-correct. If the KAG engine throws an error mid-stream, an RLVR-optimized model learns to backtrack and find a secondary reasoning path rather than collapsing.

If you want to take this further, we can look at **how KAG uses mutual indexing** to build the graph walls, or look at **how GRPO**



Medium

Knowledge Augmented...

KAG's key steps include: *
Index creation *...

GitHub

GitHub - OpenSPG/KAG:...

2.2 Mixed Reasoning Guided by Logic Forms...

GitHub

Awesome RLVR — Reinforcement...

How does it work? *
Sampling. We draw one o...

Reddit

a logical reasoning framework based c

KAG: a logical reasoning :
framework based on KG €

should this always be in a dual model system or across auditing system

August 8, 2026

No, this architecture does not always require a dual-model system, but deploying it across a **hybrid, multi-layered auditing framework** is fast becoming the industry gold standard.

While a single model can technically learn to prompt itself and call external tools, splitting the workload into dedicated roles completely eliminates critical bottlenecks in processing latency and system cost.

System Architecture Options

Frontier platforms mix RLVR and KAG across three primary architectural configurations:

1. The Monolithic System (Single Model + Extensible Tools)

- The Design:** One large model handles everything. It analyzes the user request, generates the graph-traversal plan, queries the [OpenSPG](#) knowledge base, and synthesizes the final text output.
- The Verdict: Highly inefficient.** Running an entire multi-billion parameter model through iterative self-correction loops dramatically spikes token costs and introduces severe inference latency.


2. The Dual-Model System (Generator + Internal Critic)

- The Design:** A fast, highly capable **Generator Model** drafts the logic steps. A separate, specialized **Critic Model** (trained explicitly to spot logical discrepancies) audits the generator's reasoning trail *before* passing anything to the database.
- The Verdict: Better, but still vulnerable.** This reduces latency, but because the Critic is still a neural network, it can succumb to "collusion or shared hallucinations." It might approve a broken execution path because it suffers from the same subtle bias as the Generator.

3. The Cross-Auditing System (The Industry Standard)

- The Design:** A lightweight, reasoning-optimized model produces the logic steps. These steps are then instantly pushed through an **exogenous, deterministic verification stack** that bridges neural AI with static software code.






How the Cross-Auditing Stack Functions

Instead of relying on another LLM's subjective "opinion," the Cross-Auditing method evaluates the model's performance against two completely non-neural pillars:

Audit Layer	Responsibility	Objective Verifier	Reward Mechanism
Syntactic Layer	Validates structural integrity.	Python Sandbox / JSON-Schema Linters	0.0 if the output has broken syntax or uncompileable query code.
Semantic / Graph Layer	Validates factual grounding.	KAG Mutual Indexing Graph Engine	0.0 if the path accesses non-existent entities or breaks graph constraints.

Why Cross-Auditing Defeats Dual-Model Issues

- Absolute Hard Stops:** A code compiler or database engine cannot be sweet-talked or tricked by a clever LLM. The reasoning path either strictly aligns with the enterprise schema, or it fails.
- Massive Cost Compression:** Running a deterministic Python script or graph query to verify an output costs a fraction of a penny. Running a secondary LLM to judge the output costs significantly more, making it impossible to scale during intense training runs.
- Clean Training Data:** By utilizing the environment itself as the auditor, the feedback loop generates pristine, verified synthetic training data. This cleanly drives RLVR algorithms (like GRPO) to permanently update the model's core policy.



besides semantic what other types of memory are there

August 8, 2026

Beyond semantic memory (which handles general facts and concepts), human memory is broadly divided into two major systems: **declarative (explicit)** and **non-declarative (implicit)** memory. [🔗](#)

Here is how the other types of memory break down across these categories:

1. Declarative (Explicit) Memory

This is your conscious, intentional memory for things you can explicitly cross-reference, speak aloud, or recall. Along with semantic memory, it includes: [🔗](#)

- **Episodic Memory:** Your personal mental diary. It stores unique, autobiographical experiences tied to a specific time and place (e.g., your first day at a new job or what you ate for breakfast today). [🔗](#)

2. Non-Declarative (Implicit) Memory

This system functions automatically without conscious awareness. You do not intentionally "retrieve" these memories; instead, they influence your actions and behaviors. [🔗](#)

- **Procedural Memory:** Your physical and motor skills blueprint. It is how your brain remembers *how* to do things automatically (e.g., riding a bicycle, typing on a keyboard, or playing an instrument). [🔗](#)
- **Priming:** The phenomenon where exposure to one stimulus influences your subconscious response to a subsequent stimulus (e.g., if you recently saw the word "doctor," you will recognize the word "nurse" significantly faster). [🔗](#)
- **Classical Conditioning:** The automatic association made between a stimulus and a behavioral response, like flinching when you hear a sound previously paired with a loud pop. [🔗](#)

3. Temporal Categories (Duration-Based)

Memory can also be classified by how long the information is held before it is processed or discarded: [🔗](#)

- **Sensory Memory:** The ultra-short-term holding tank. It retains raw, literal environmental inputs for a fraction of a second (e.g., iconic memory for visual snapshots, echoic memory for auditory sounds). [🔗](#)
- **Working Memory:** The active mental workspace. It temporarily holds and actively manipulates a limited amount of information for immediate task execution (e.g., holding a phone number in your head just long enough to type it out). [🔗](#)
- **Short-Term Memory:** The passive storage of small amounts of information for a brief duration (around 15–30 seconds) without active manipulation. [🔗](#)



- Would you like to look at **how these human systems map to AI architectures** (like vector embeddings vs. long-term weights)?
- Should we explore the **neurology behind how episodic memory shifts into semantic memory** over time?



Pressbooks.pub

8.2: Types of Memory –...

Nondeclarative Memory (Implicit) The other main...

Study.com

Types of Memory | Overview &...

Explicit Memory Explicit memory is generally...

Internet Encyclopedia of Philo...

Memory, Epistemology of

At least as far back as Henri Bergson (1896/1994) and Bertrand Russell...

ResearchGate

(PDF) Models of semantic memory

Semantic Memory: Memory for word...

hey what's that Gina chat or Gina something

August 8, 2026

You are thinking of [Jina AI](#), a major developer platform specialized in search, retrieval, and AI data engineering. While people sometimes look for a "Jina chat" interface, Jina is actually famous for providing the infrastructure behind chat systems, AI agents, and RAG (Retrieval-Augmented Generation). YouTube · AI An... +1

The primary tools that have made Jina popular across the AI ecosystem include:

1. Jina Reader API

This is their most famous utility for AI developers. If you append `r.jina.ai/` to the front of any web link (e.g., `r.jina.ai/https://wikipedia.org`), it instantly strips away ads, tracking scripts, popups, and raw HTML. It outputs pristine, clean **markdown text**. Developers feed these outputs into chat models so the AI can read real-time web content without choking on junk code. YouTube · AI An... +1

2. High-Performance Embeddings & Rerankers

Jina builds specialized, high-context models that convert text, code, or images into mathematical vectors. This tech acts as the long-term semantic memory for LLMs, allowing enterprise chatbots to fetch exact answers from thousands of page documents in milliseconds. LinkedIn · Jina AI +2

3. DeepSearch & DeepResearch

Jina offers specialized APIs designed to empower autonomous AI agents to execute deep, multi-layered web searches and research pipelines. They use multilingual architectures to cross-reference and rank sources intelligently before presenting a synthesis to an end-user.

Where You Might Have Encountered "Jina Chat"

- **Third-Party Chat UIs:** Because Jina offers an API key, many users plug their Jina credentials into wrapper platforms like [TypingMind](#) to run customized web-search chats. TypingMind



YouTube · In The Kitchen With ...

GINA YOUNG LIVE CHAT PULL UP

GINA YOUNG LIVE CHAT PULL UP Hello Everyone...

YouTube · AI Anytime

Jina Reader API: Build better AI...

so Gina AI has created reader. they have come u...

LinkedIn

Jina AI | LinkedIn

Founded by Dr. Han Xiao in 2020, Jina AI is a leading...

Google Cloud

How Jina AI built it 100-billion-token..

Editor's note: The Jina AI Reader is a specialized...

yeah that's it generator what about the one where you can put PDFs in and get txt

▼

August 8, 2026

Yes, you are thinking of the [Jina Reader API](#). It handles PDFs natively and extracts all the text. [🔗](#)

If you have a PDF hosted on the web, you use it exactly like a normal webpage: just append `r.jina.ai/` to the front of the PDF's URL. [🔗 Elastic — The S... +1](#)

How to Use the Jina PDF Feature

- **For Web PDFs:** Paste the URL into your browser like this:
`r.jina.ai/https://example.com`
- **The Output:** It strips away the complex PDF layout and instantly hands you clean, structured **Markdown text**.
- **Image Support:** It uses built-in vision models, meaning if your PDF contains charts, complex tables, or images, it will actually generate **text descriptions** of those images directly in the file. [🔗](#)

For Local PDFs (Uploaded from your computer)

If you aren't pulling a PDF from a website and want to upload your own files, developers use **Jina's API endpoints** (often integrated into Chat interfaces like [TypingMind](#) or automation tools like [Relevance AI](#)) to process raw file uploads into LLM-friendly text streams. [🔗 ColdIQ +1](#)

Are you hoping to **convert a massive batch of local PDFs**, or are you looking for a **no-code web interface** where you can just drag-and-drop a file to get text back?



Elastic — The Search AI Comp...

Jina Reader - Elasticsearch Labs

Elasticsearch has a service called Jina Reader that...

X

Jina AI on X: "PDFs are hard knuckle,..."

PDFs are hard knuckle, and you know it. Reader...

DEV DEV Community

markdown.new + Jina Reader: Stop...

Rubbish way to spend your tokens. markdown.new...

Jina AI

Reader API - Jina AI

Reader also reads PDFs! Yes, Reader natively...

🔊

Yes, you can do this completely without an API or writing code by using the graphical interface on the [Jina Reader Web Dashboard](#). 

The Easiest Way to Convert a Local PDF (No-Code)

1. **Open the Webpage:** Go directly to the official **Jina Reader Page** in your web browser.
2. **Locate the Uploader:** Scroll slightly down the page to find the configuration dashboard.
3. **Upload the File:** Look for the option labeled **Local PDF/HTML file**. There is a direct **Upload button** right next to it.
4. **Get Your Text:** Drag and drop your single PDF document from your computer. The web tool will process it using their system right in your browser window and return the clean, extracted text instantly. 

Alternative: The Free Google Drive Trick

If you ever find Jina's web dashboard busy or restricted, you can get clean text from a single PDF using standard cloud tools you likely already have:  [Reddit](#) +1

1. Upload your PDF to your **Google Drive**.
2. Right-click the PDF file, select **Open with**, and choose **Google Docs**.
3. Google will instantly run its own optical character recognition (OCR) loop, stripping away the layout formatting and opening the raw text in a fresh document.

Would you like help looking at **other browser-based document extraction tools**, or are you ready to try **uploading your file directly to Jina**?  Jina AI



 Jina AI
Reader API - Jina AI
Include Shadow DOM.
Extract content from...

 Reddit
Best way to extract text from PDF docs for finetuning...
Apache Tika (via Java, or commandline JAR) does an excellent job of extractin...

 Reddit
Extracting text from a PDF without using PyPDF2 :...
According to a post on r/learnpython, there are a few ways to extract text fro...

 www.lido.app
How to Extract Text from a PDF (5 Methods) -
July 1, 2026 In this article scanned PDFs: why it ma

nope that one wanted to charge me

August 8, 2026


Ah, apologies for that! Their dashboard can definitely sneak a credit check or prompt you to buy token bundles if you try to drag-and-drop directly into their advanced console. 

If you just have one PDF and want a completely free, zero-signup web interface to extract the text, use [OnlineOCR](#) or [Convertio PDF to TXT](#).



- **OnlineOCR:** Ideal if your PDF contains scanned images or photos of text. It runs a clean character-recognition loop for free right on the homepage and spits out the text in a copy-pasteable box.
- **The Microsoft Edge Sidebar Trick:** If you use Windows, open the PDF right inside the Microsoft Edge browser. Click the **Copilot** icon in the top right corner and type: *"Extract all the text from this page and give it to me as markdown."* It will read the open document instantly for free.

If you are looking to pull text out of a file that has **heavy charts, math equations, or tables**, let me know so I can point you to a tool that won't scramble the layout!



you got a PDF to txt converter that actually works dude

August 8, 2026

[PDF24 Tools](#) or [iLovePDF](#) are completely free web tools that actually work for single document conversions without forcing you to pay, upgrade, or sign up for an API key.

If you are dealing with text extraction, you have three primary "zero-dollar, zero-signup" options depending on where your document is sitting:

1. The Pure Web Browser Route (No Accounts, Free)

- **For Normal PDFs:** Drop your file directly into the [PDF24 PDF-to-Text Converter](#) or [Drawboard PDF to TXT](#). They process the layout instantly and spit back a downloadable `.txt` file. Drawboard P... +1
- **For Scanned Photos / Images of Text:** Use the [iLovePDF OCR Engine](#). It reads the flattened text embedded inside the image lines and converts it into selectable text blocks. iLovePDF

2. The Local Machine Desktop Workarounds

If you don't want to upload sensitive data to an internet server, use software you already have installed:

- **Microsoft Word:** Open Microsoft Word, click **File > Open**, and select your PDF. Word will run a local extraction loop to strip the formatting and turn it into editable text. You can then save it right back down as a plain text `.txt` file.
- **The Microsoft Edge Sidebar:** Open your PDF file inside the Edge browser. Click the **Copilot** sidebar panel and type: *"Copy every piece of text from this open document into a markdown snippet."*

3. The 3-Line Python Script (100% Free & Unlimited)



```
pip install pymupdf
```

Use code with caution. 

Then, execute this tiny block of code to generate your text file: 

```
python

import pymupdf

# Open your document
doc = pymupdf.open("your_file.pdf")

# Extract text and save it straight to a clean TXT file
with open("extracted_output.txt", "w", encoding="utf-8") as f:
    for page in doc:
        f.write(page.get_text())
```

Use code with caution. 

Let me know what you'd like to do next:

- If you want to use the script, I can show you **how to run it locally** step-by-step.
- If you have a specific file layout (like **multi-column sheets or financial tables**), I can point you to the tool that preserves that alignment.





PyMuPDF is a high performance...

Why PyMuPDF? * Fast — powered by MuPDF, a...





any actual free pdf to text converters...

According to a Reddit user, the only free PDF to text...





Best PDF text extractor in Pytho...

Extracting the complete text from PDF for further...





OCR PDF to text tool: Make PDF...

How to make a PDF searchable. Need to kno..

GitHub

August 12, 2026

Quick results from the web:

GitHub · Change is constant. GitHub keeps you ahead. · GitHub

From your first line of code to final deployment, GitHub provides AI and automation tools to help you build and ship better software faster. level roadmaps to ...

 GitHub · <https://aithub.com> 



If you want to use the **GitHub** route to convert your PDF, the easiest way is to use a free, browser-based app hosted right on GitHub Pages. This runs completely inside your web browser. Your file is never uploaded to a server, making it 100% free, private, and secure. 

1. The 1-Click Browser Tool (Hosted on GitHub)

You can use **PDF to Text by OCR.space on GitHub** or [Stirling-PDF Web](#). 

- **How to use it:** Open the page, drag your PDF into the box, and click convert.
- **Why it works:** It uses your browser's local power to read the text and save it as a `.txt` file instantly. 

2. Ready-to-Run GitHub Repositories

If you want to run a dedicated tool on your computer, developers have built incredible free tools that you can download directly from GitHub: 

- [Stirling-PDF](#): This is the ultimate open-source PDF toolkit. It gives you a gorgeous web interface on your local computer to convert, merge, split, and run OCR on PDFs for free forever. 
- [Marker](#): If your PDF has complex math formulas, charts, or multiple columns, this GitHub project converts it into perfectly clean Markdown text. 
- [Nougat](#): Built by Meta AI, this tool is specifically designed on GitHub to extract text from scientific PDFs and academic papers without scrambling the layout.


Would you like help **setting up one of these GitHub tools** on your computer, or would you prefer a link to a **live GitHub Pages demo** you can use right now?



 Hackney Council

Using GitHub | Hackney Data...

Getting access to GitHub Using the GitHub web...



 DAEXT

Create a Personal Website With...

If you're looking for the easiest possible route,...



 Reddit

When Should I Learn Git/GitHub? ...

GitHub isn't something you need to learn, but it's probably the easiest way...

 Python in Plain English

From Zero to Deployed: Buildin..

The best part about this approach is that you can.