To synthesize, de-duplicate, and organize these overlapping architectural frameworks, tracking files, and researcher data banks into a clean, unified structure for unilateral analysis, we need to categorize them by operational layer.
Looking at your files, your ecosystem has evolved through multiple naming conventions and target iterations—ranging from early terminal-native sandboxes (**AESOP**, **Termux CLI Stack**) to native on-device hardware engines (**Novus-Agenti**, **NovA-Claw**, **N0.V4**, and the upgraded **Horizons Workbench** UI).
The canonical repository blueprint requires sorting everything into **Four Core Buckets** to prevent your engineering agents from drifting into old contexts or hallucinating stale runtime targets.

Tier 1: The Cold Archive & Source References (/raw_sources/)

These are heavy, static background documents, technical metrics, and community tutorials that the model should use strictly as read-only reference data—never as active working rules.

- **Stanford AI Index & General Trends:**
  - Artificial Intelligence Index Report | Stanford HAI (Massive industry benchmark data bank)
- **External Technical / Framework Discoveries:**
  - MEMO_ A Modular Framework for Training a Dedicated Memory Model...pdf (Parametric memory layer research)
  - Together AI Open-Sources OSCAR...pdf (2-bit attention-aware KV cache optimization spec)
  - NVIDIA Releases Polar...pdf (Gateway rollout framework for token-faithful GRPO training loops)
  - Meet GitAgent..., OpenAI Releases Symphony..., How CopilotKit Is Redefining..., Meet GitHub Spec-Kit... (Framework mappings for agentic orchestration)
  - Meet OmniVoice Studio...pdf (Tauri/FastAPI local alternative reference stack)
  - Hexo Labs Open-Sources SIA... (Self-improving agent scaffold loops)
- **Qualcomm / Hardware Specific Specs:**
  - qualcomm/Phi-4-Mini-Instruct · Hugging Face, Phi-4-Mini-Instruct - Qualcomm AI Hub
  - qualcomm/Qwen3.5-0.8B · Hugging Face, Qwen3.5-0.8B - Qualcomm AI Hub, unsloth/Qwen3.5-9B-GGUF · Hugging Face
  - Gemma-4-E4B-it - Qualcomm AI Hub
  - qualcomm/LiteHRNet · Hugging Face, qualcomm/PiperTTS-DE · Hugging Face
  - Liquid AI Releases LFM2.5-8B-A1B...
  - 8.2mb what is a hybridized agentic on device neuro mesh...pdf
- **Keyboard & Input Utilities:**
  - FUTO Keyboard, Voice Input Models - FUTO Keyboard

Tier 2: Core Architecture Vault (/master_wiki/)

This is the single source of truth for design constraints, structural layouts, and evolution reports mapping out the native system topology.

- **Native Android (Kotlin/Compose) Next-Gen Workbenches:**
  - HORIZONS_state_and_corrections_MASTER.md (**Canonical State Document** detailing

structural panel layouts, model layer bindings, and the credential list rewrite)
    - ARCHITECTURE_REVIEW_v3_amendment.md (Crucial dual-axis framework detailing the Cloud Failover Ladder vs. the VLM Context Dispatcher)
    - The Horizons Workbench: A Guide to Your Modular Intelligence Engine
    - 4. Technical Specification: The Four Rooms & Seven-Tile Modular Architecture
    - The Horizons Workflow: From Storage to Switch
    - HORIZONS_UI_SPEC_v3.md (Detailed single-Activity layout and Accessibility Service automation blueprints)
    - N0_V4_ARCHITECTURE_v3.md (Distributed edge-cloud specification targeting Snapdragon 8 Elite hardware metrics)
    - Architecting the Agentic Neural Mesh for Snapdragon 8 Elite, Agentic Neural Mesh: Master Architecture and Evolution Report
- **Legacy / Interim Sandbox Ecosystem Templates (Marked as Contextual Reference):**
    - TERMUX_CLI_STACK.md (Bridges the gap; maps out the native headless terminal environment layer utilized by RUN_COMMAND)
    - Copy of SOTU-2026-07-27.md (Historical project state tracking across legacy variations)
    - deploy_T3-BRINGUP.md (Infrastructure specifications mapping the Tier 3 home network hub target: Jetson Orin Nano Super + Rubik Pi 3 desktop-OS node over Tailscale)

Tier 3: The Almanac & Skill Definitions (/inference_skills/)

Line-delimited facts, structured JSONL logs, terminal execution setups, and open SKILL.md plugin files used to instantly direct local engines.

- **Active Agent Tool Routing Rules:**
    - horizons-wiki-SKILL.md (Packages stable code-review architecture, daemon path parameters, and Anthropic ephemeral caching boundaries into a system block prefix)
    - 1b How to Structure Your Wiki Folders..., Turning a document into a Skill by adding Frontmatter... (Operational instructions for compiling modular .txt or .json directories into queryable skill blocks)
    - claude-skills.md (Reference domain bundle for external tool building/auditing frameworks)
- **Terminal Environment & API Wiring Recipes:**
    - GENIEX-DAEMON-PLAN.md / Copy of GENIEX-DAEMON-PLAN.md (Blueprints for launching standard OpenAI request formats via a standalone geniex serve native process over port 18181)
    - Standardized UNIX socket protocol, (Zero-copy IPC optimization blueprints utilizing shared memory fields)
    - llm-wiki-termux-setup-full-notes.md / llm-wiki-termux-setup-urls.txt (Headless environment installation logic)

Tier 4: Session Execution & Hygiene Logs (/automation_scripts/)

Active working tracks, debugging sequences, failure logs, and step-by-step master checklists used to eliminate sessional drift.

- **Current Iteration Logs:**
    - BUILD-ACTION-PLAN.md / Copy of BUILD-ACTION-PLAN.md (Parallel processing

partitioning guidelines for Repo vs. Device workflows)
    - FAILURE_LOG_SESSION_20260527.md / weekly_audit_prompt.md (Failure frequency tracker monitoring degradation types like specification drift or tool error loops)
    - NeuroOmni_agy_build_pack.md (Milestone execution checkpoints for scaffolding the base project architecture)

Step-by-Step Execution for Unified Analysis

To prevent context corruption, do not load heavy references simultaneously with the active target code. Execute your unilateral analysis using the following pipeline:

1. **Lock the Base Pool:** Use **HORIZONS_state_and_corrections_MASTER.md** combined with **ARCHITECTURE_REVIEW_v3_amendment.md** as your absolute architectural anchor.
2. **Evaluate Path Congruence:** Use **BUILD-ACTION-PLAN.md** to verify that file alterations stay completely distinct across active sub-sessions.
3. **Run the Audit Matrix:** Filter incoming tracking data using the failure categories detailed in **weekly_audit_prompt.md** to verify execution readiness before generating patches.